

Rohan Jha

rjha@andrew.cmu.edu • robro612@gmail.com • 858-229-3216

San Diego, California • Pittsburgh, Pennsylvania

EDUCATION

Carnegie Mellon University

B.S. Artificial Intelligence

May, 2023

Pittsburgh, PA

- GPA: 3.89, Dean's List High Honors
- Relevant Coursework: Parallel and Sequential DS and Algorithms (15-210), Introduction to Machine Learning (10-315), Search Engines (11-442), Introduction to Deep Learning (11-485), Convex Optimization (10-725)

SKILLS & INTERESTS

- **Languages:** Python, C, Java, Standard ML, Bash
- **Technologies:** Git, Mercurial, NumPy, PyTorch, LaTeX, Junit & Mockito, AWS, Slurm, Conda, W&B
- **Concepts:** Optimization, Machine Learning, Data Structures & Algorithms, Functional Software Development
- **Interests:** Information Retrieval, Representation Learning, AI Robustness & Explainability

EXPERIENCE

Software Engineering Intern

Meta (Ads Core ML Modeling Intelligence)

May, 2022 – August, 2022

Menlo Park, CA

- Using Caffe2/PyTorch, implemented sparse Mixture of Experts and novel DSelectK Gating techniques into multiple sub-architectures of production advertisement recommendation models
- Conducted validation and ablation experiments to determine the efficacy and infrastructure costs of newly introduced modules, achieving model performance improvements supporting multiple organizations across the company corresponding to increased revenue

Research Assistant

Carnegie Mellon University (Language Technology Institute)

December, 2021 – Present

Pittsburgh, PA

- Supported grad-level research in neural information retrieval focused on combining dense language models' context with the sparse efficiency of the inverted list architecture
- Designed, coded, and presented experiments and results to the principal investigator
- Produced conference-style work analyzing the performance/cost tradeoff and distribution of sparsely-factorized dense embeddings as a performance-preserving, cost-optimizing modification to an existing retriever

Teaching Assistant

Carnegie Mellon University (15-281: AI: Representation and Problem Solving, 07-180: Concepts of AI)

January, 2021 – May, 2022

Pittsburgh, PA

- Designed, tested, proctored, and graded written and programming-based homework and exams
- Led 120+ students with weekly office hours, recitations, and exam review sessions
- Was primary expert for student questions on class forum with high coverage and low response latency

Software Development Engineer Intern

Amazon (Prime Video Ads)

June, 2021 – August, 2021

Seattle, WA

- Designed, implemented, and tested software processing billions of customer impression records each hour
- Led extensive design and review process after researching and documenting various solution alternatives
- Reduced latency, cost, and on-call pain points to accommodate imminent 3-7x increase in impression volume