

Rohan Jha

rjha@utexas.edu • robro612@gmail.com • 858-229-3216
San Diego, California • Austin, Texas

EDUCATION

The University of Texas at Austin

M.S. Computer Science

May, 2025

Austin, TX

Carnegie Mellon University

B.S. Artificial Intelligence

May, 2023

Pittsburgh, PA

- GPA: 3.91, Dean's List High Honors, Phi Beta Kappa
- Relevant Coursework: ML with Large Datasets, Search Engines, Question Answering, Deep RL

SKILLS & INTERESTS

- **Languages:** Python, C, Java, Standard ML, Bash
- **Technologies:** Git, Mercurial, NumPy, PyTorch, LaTeX, Junit & Mockito, AWS, Slurm, Conda, W&B
- **Concepts:** Optimization, Machine Learning, Data Structures & Algorithms, Functional Software Development
- **Interests:** Information Retrieval, Representation Learning, Question Answering, Machine Learning

PUBLICATIONS

- "COILCR: Efficient Semantic Matching in Contextualized Exact Match Retrieval"
Zhen Fan, Luyu Gao, **Rohan Jha**, Jamie Callan.
In *Advances in Information Retrieval (ECIR)*, 2023.

EXPERIENCE

Software Engineering Intern

Meta (Ads Core ML Modeling Intelligence)

May, 2022 – August, 2022

Menlo Park, CA

- Using Caffe2/PyTorch frameworks, implemented sparse Mixture of Experts and novel DSelectK Gating techniques into multiple sub-architectures of production advertisement recommendation models
- Conducted validation and ablation experiments to determine the efficacy and infrastructure costs of newly introduced modules, achieving model performance improvements supporting multiple organizations across the company corresponding to increased advertisement revenue

Research Assistant

Carnegie Mellon University (Language Technology Institute)

December, 2021 – December, 2022

Pittsburgh, PA

- Supported published research in neural information retrieval focused on combining dense language models' context with the sparse efficiency of the inverted list architecture
- Designed, implemented, and presented experiments and results to the principal investigator
- Produced conference-style work analyzing the performance/cost tradeoff and distribution of sparsely-factorized dense embeddings as a performance-preserving, cost-optimizing modification to an existing retriever

Teaching Assistant

Carnegie Mellon University (15-281: AI: Representation and Problem Solving, 07-180: Concepts of AI)

January, 2021 – May, 2022

Pittsburgh, PA

- Designed, tested, proctored, and graded written and programming-based homework and exams
- Was primary expert for questions from 120+ students on class forum with high coverage and low response latency